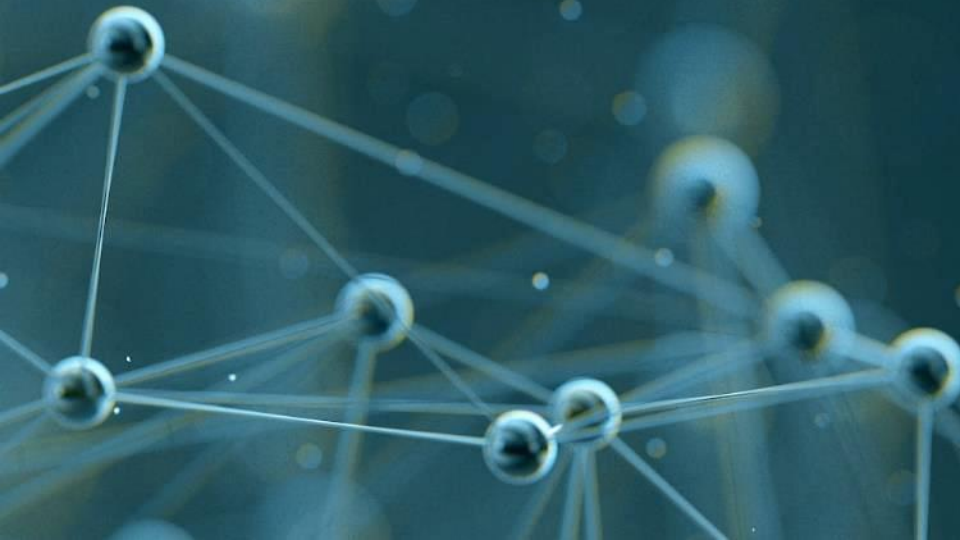



Support Vector Machines




Relationship to Logistic Regression




$$y_{\beta}(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$

Support Vector Machines (SVM)




Support Vector Machines (SVM)




Three misclassifications

Support Vector Machines (SVM)




Two misclassifications

Support Vector Machines (SVM)




No misclassifications

Support Vector Machines (SVM)




No misclassifications—but is this the best position?

Support Vector Machines (SVM)




No misclassifications—but is this the best position?

Support Vector Machines (SVM)



Maximize the region between classes.


Similarity Between Logistic Regression and SVM



Classification with SVMs


Two features (nodes, age)

Two labels (survived, lost)




Classification with SVMs

Find the line that best separates classes.




Classification with SVMs

Find the line that best separates classes.




Classification with SVMs

Find the line that best separates classes.




Classification with SVMs

Find the line that best separates classes.




Classification with SVMs


Also, include the largest boundary possible.





Logistic Regression vs SVM Cost Functions




Logistic Regression vs SVM Cost Functions




Logistic Regression vs SVM Cost Functions





The SVM Cost Function





The SVM Cost Function




The SVM Cost Function




The SVM Cost Function




The SVM Cost Function




Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs



Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs

This is probably still the correct boundary.




Regularization in SVMs

$$J(\beta_i) = \text{SVMCost}(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



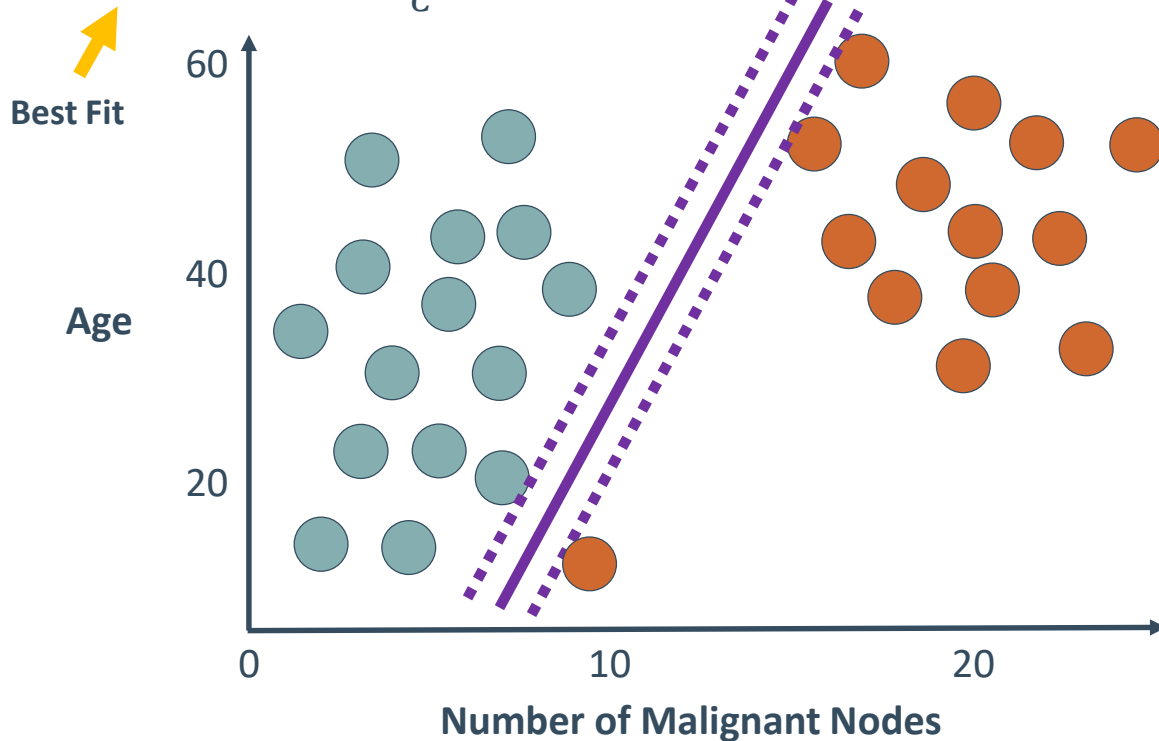
Regularization in SVMs

$$J(\beta_i) = \text{SVMCost}(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Regularization in SVMs


$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Regularization in SVMs


$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$

← Large




Regularization in SVMs

$$J(\beta_i) = \text{SVMCost}(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$




Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$


← Much Smaller

↑
Slightly Higher




Interpretation of SVM Coefficients

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Interpretation of SVM Coefficients

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Linear SVM: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import LinearSVC
```

Linear SVM: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import LinearSVC
```

Create an instance of the class.

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```


Linear SVM: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import LinearSVC
```

Create an instance of the class.

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```



Regularization
parameters

Linear SVM: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import LinearSVC
```

Create an instance of the class.

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```

Fit the instance on the data and then predict the expected value.

```
LinSVC = LinSVC.fit(X_train, y_train)  
y_predict = LinSVC.predict(X_test)
```

Linear SVM: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import LinearSVC
```

Create an instance of the class.


```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```

Fit the instance on the data and then predict the expected value.


```
LinSVC = LinSVC.fit(X_train, y_train)  
y_predict = LinSVC.predict(X_test)
```

Tune regularization parameters with cross-validation.

Kernels




Classification with SVMs




Non-Linear Decision Boundaries with SVM

Non-linear data can be made linear with higher dimensionality.




The Kernel Trick

Transform data so it is linearly separable.




SVM Gaussian Kernel

Palme d'Or Winners at Cannes.



SVM Gaussian Kernel

Palme d'Or Winners at Cannes.




Approach 1:
Create higher order
features to transform
the data.

$\text{Budget}^2 +$
 $\text{Rating}^2 +$
 $\text{Budget} * \text{Rating} +$
...

SVM Gaussian Kernel


Palme d'Or Winners at Cannes.



Approach 2:
Transform the space to
a different
coordinate system.

SVM Gaussian Kernel


Palme d'Or Winners at Cannes.



Define Feature 1:
Similarity to
“Pulp Fiction.”

SVM Gaussian Kernel


Palme d'Or Winners at Cannes.



Define Feature 2:
Similarity to
"Black Swan."

SVM Gaussian Kernel

Palme d'Or Winners at Cannes.




Define Feature 3:
Similarity to
"Transformers."

SVM Gaussian Kernel

Palme d'Or Winners at Cannes.

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Pulp Fiction})^2}{2\sigma^2} \right]$$




Create a
Gaussian function
at feature 1.

SVM Gaussian Kernel

Palme d'Or Winners at Cannes.

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Black\ Swan})^2}{2\sigma^2} \right]$$




Create a
Gaussian function
at feature 2.

SVM Gaussian Kernel


Palme d'Or Winners at Cannes.

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Transformers})^2}{2\sigma^2} \right]$$




Create a
Gaussian function
at feature 3.


SVM Gaussian Kernel



SVM Gaussian Kernel




SVM Gaussian Kernel



SVM Gaussian Kernel

Transformation:


$$[x_1, x_2] \rightarrow [0.7a_1, 0.9a_2, -0.6a_3]$$



Classification in the New Space


Transformation:

$$[x_1, x_2] \rightarrow [0.7a_1, 0.9a_2, -0.6a_3]$$




SVM Gaussian Kernel

Palme d'Or Winners at Cannes.




SVM Gaussian Kernel

Palme d'Or Winners at Cannes.



SVM Gaussian Kernel

Palme d'Or Winners at Cannes.



SVMs with Kernels: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import SVC
```

SVMs with Kernels: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import SVC
```

Create an instance of the class.

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

SVMs with Kernels: The Syntax


Import the class containing the classification method.

```
from sklearn.svm import SVC
```

Create an instance of the class.

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

Set kernel and
associated coefficient
(gamma).



SVMs with Kernels: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import SVC
```

Create an instance of the class.

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

"C" is penalty
associated with
the error term.



SVMs with Kernels: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import SVC
```

Create an instance of the class.

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

Fit the instance on the data and then predict the expected value.

```
rbfSVC = rbfSVC.fit(X_train, y_train)  
y_predict = rbfSVC.predict(X_test)
```

SVMs with Kernels: The Syntax

Import the class containing the classification method.

```
from sklearn.svm import SVC
```

Create an instance of the class.

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

Fit the instance on the data and then predict the expected value.

```
rbfSVC = rbfSVC.fit(X_train, y_train)  
y_predict = rbfSVC.predict(X_test)
```

Tune kernel and associated parameters with cross-validation.

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data.

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data.

Solution

- Construct approximate kernel map with SGD using Nystroem or RBF sampler

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data.

Solution

- Construct approximate kernel map with SGD using Nystroem or RBF sampler
- Fit a linear classifier

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class.

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```

Fit the instance on the data and transform.

```
X_train = nystroemSVC.fit_transform(X_train)  
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class.

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```



Multiple non-linear
kernels can be used.

Fit the instance on the data and transform.

```
X_train = nystroemSVC.fit_transform(X_train)  
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class.

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```



Kernel and gamma are identical to SVC.

Fit the instance on the data and transform.

```
X_train = nystroemSVC.fit_transform(X_train)  
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class.

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```



n_components is
number of samples.

Fit the instance on the data and transform.

```
X_train = nystroemSVC.fit_transform(X_train)  
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class.

```
rbfSample = RBFsampler(gamma=1.0,  
                       n_components=100)
```

Fit the instance on the data and transform.

```
X_train = rbfSample.fit_transform(X_train)  
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class.

```
rbfSample = RBFsampler(gamma=1.0,  
                        n_components=100)
```



RBF is only kernel that
can be used.

Fit the instance on the data and transform.

```
X_train = rbfSample.fit_transform(X_train)  
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method.

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class.

```
rbfSample = RBFsampler(gamma=1.0,  
                       n_components=100)
```



Parameter names
are identical
to previous.

Fit the instance on the data and transform.

```
X_train = rbfSample.fit_transform(X_train)  
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

When to Use Logistic Regression vs SVC

Features	Data	Model Choice
Many (~10K Features)	Small (1K rows)	Simple, Logistic or LinearSVC
Few (<100 Features)	Medium (~10k rows)	SVC with RBF
Few (<100 Features)	Many (>100K Points)	Add features, Logistic, LinearSVC or Kernel Approx.

